

NL2SQL Agent for Mako Group

Overview

The NL2SQL Agent is a natural-language-to-SQL system built for Mako Group's options market-making desk. Traders ask questions in plain English -- "what was the edge on our quoter trades yesterday?" -- and the agent routes to the correct BigQuery table, generates SQL, validates it, executes it, and returns the answer.

It runs as a **Google ADK sub-agent** behind a LiteLLM proxy, consumed by CLI agents (Claude Code in dev, Gemini CLI in production).

Architecture

```
Trader: "what was PnL by delta bucket yesterday?"
|
v
Root Agent (mako_assistant)
|   Delegates data questions
v
NL2SQL Sub-Agent (nl2sql_agent)
|
|-- 1. vector_search_tables      --> Semantic search: which
table(s)?
|-- 2. load_yaml_metadata       --> Column names, types,
synonyms, rules
|-- 3. fetch_few_shot_examples --> Similar validated Q->SQL
pairs
|-- 4. [LLM generates SQL]
|-- 5. dry_run_sql               --> Validate syntax, estimate
cost
|-- 6. execute_sql              --> Run query, return results
|-- 7. save_validated_query     --> Save for future retrieval
(learning loop)
|
v
BigQuery Results --> Formatted answer to trader
```

Agent Structure

Agent	Role	Model	-----	-----	-----	mako_assistant (root)	Routes questions. Data questions -> nl2sql_agent. General questions -> direct answer.	LiteLLM proxy	nl2sql_agent (sub-agent)	6-tool chain for SQL generation and execution. Temperature 0.1.	LiteLLM proxy
-------	------	-------	-------	-------	-------	-----------------------	---	---------------	--------------------------	---	---------------

Tech Stack

Component	Technology	-----	-----	Agent Framework	Google ADK v1.25.1	LLM Access									
LiteLLM proxy	(Claude in dev, Gemini in prod)	Database	BigQuery (read-only)	Embeddings	text-embedding-005 via BigQuery ML	Vector Search	BigQuery VECTOR_SEARCH (COSINE, brute-force)	Config	pydantic-settings + .env	Logging	structlog (JSON)	Container	Docker + docker-compose	Testing	pytest (155 unit + 20 integration tests)

Data Model: 13 Tables Across Two Datasets

KPI Dataset (n12sql_omx_kpi) -- Gold Layer

Performance metrics per trade. One table per trade origin, all sharing KPI columns (edge_bps, instant_pnl, delta_bucket, slippage intervals).

| Table | What It Contains | When to Use | -----|-----|-----| | **markettrade** | Exchange trade KPIs | **DEFAULT** when trade type unspecified || **quotertrade** | Auto-quoter fill KPIs | "quoter edge", "quoter PnL" || **brokertrade** | Broker trades (has account field) | "BGC vs MGN", broker comparison || **clicktrade** | Manual click trade KPIs | "manual trade performance" || **otoswing** | OTO swing trade KPIs | OTO-specific metrics |

Data Dataset (n12sql_omx_data) -- Silver Layer

Raw execution details, timestamps, prices, market data.

| Table | What It Contains | When to Use | -----|-----|-----| | **theodata** | Theo pricing: vol, delta, vega, theta | "implied vol", "greeks", "fair value" || **quotertrade** | Raw quoter activity (timestamps, levels) | "what levels at 11:15?" || **markettrade** | Raw trade execution details | Raw timestamps, fill prices || **clicktrade** | Raw click execution details | Raw click activity || **marketdata** | Market price feeds (top-of-book) | "market price", "price feed" || **marketdepth** | Order book depth, multiple levels | "order book", "bid-ask levels" || **swingdata** | Raw swing trade data | "swing latency" |

Critical Routing Challenge

kpi.quotertrade and data.quotertrade are **different tables** with similar names. The agent must disambiguate based on question intent:

- "What was quoter edge today?" -> kpi.quotertrade (performance metrics)
- "What levels were we quoting at 11:15?" -> data.quotertrade (raw activity)

Same ambiguity exists for markettrade and clicktrade.

Two-Layer Metadata System

The agent uses two complementary layers to understand what data exists and how to query it.

Layer 1: YAML Catalog (Static, In-Repo)

One YAML file per table stored in catalog/. Provides:

- **Column descriptions:** what each column means in trading context
- **Synonyms:** maps trader language to column names (e.g. "edge" -> edge_bps)
- **Business rules:** how calculations work, what values mean
- **Routing rules:** which table to use for which question type
- **Disambiguation:** resolves overlapping table names across datasets

```
catalog/
    _routing.yaml          # Cross-dataset routing rules
    kpi/
        _dataset.yaml      # KPI shared columns + routing
    patterns
        markettrade.yaml   # 150+ columns with descriptions
        quotertrade.yaml
```

```

brokertrade.yaml
clicktrade.yaml
otoswing.yaml
data/
    _dataset.yaml           # Data dataset context
    theodata.yaml
    quotertrade.yaml
    markettrade.yaml
    clicktrade.yaml
    marketdata.yaml
    marketdepth.yaml
    swingdata.yaml

```

Example (kpi/markettrade.yaml excerpt):

```

table:
    name: markettrade
    dataset: nl2sql_omx_kpi
    fqn: "{project}.nl2sql_omx_kpi.markettrade"
    partition_field: trade_date
    columns:
        - name: edge_bps
          type: FLOAT64
          description: >
              Edge captured in basis points. Positive = better than fair
              value.
              Primary performance metric for traders.
              synonyms: [edge, the edge, how much edge, trading edge]
        - name: instant_pnl
          type: FLOAT64
          description: Immediate profit/loss at execution, native
                      currency
          synonyms: [PnL, profit, loss, p&l]

```

The {project} placeholder is resolved at runtime from settings, allowing the same catalog to work across dev and prod environments.

Layer 2: BigQuery Vector Embeddings (Dynamic, In BQ)

Three tables in the nl2sql_metadata dataset, all using text-embedding-005 (768-dim vectors):

Table	Rows	Purpose	Used By	----- ----- ----- -----	schema_embeddings	~17	Dataset and table descriptions
vector_search_tables					column_embeddings	~1000	Column names + descriptions + synonyms
Future column routing					query_memory	30+ (growing)	Validated question->SQL pairs fetch_few_shot_examples

Why Two Layers?

Concern	YAML Catalog	Vector Embeddings	----- ----- -----	Column detail	Full descriptions, types, synonyms Not needed	Semantic search	Not searchable COSINE similarity
Maintenance					Edit YAML, commit to git Auto-populated from YAML + scripts	Learning	Static query_memory grows via save_validated_query
					Offline	Works without BQ Requires live BQ connection	

The YAML catalog is the **source of truth** for metadata. The embeddings are a **search index** over that metadata, plus a growing memory of validated queries.

How Embedding and Vector Search Works

Embedding Pipeline

```
1. YAML catalog files .....> populate_embeddings.py .....> BigQuery  
tables  
    (column descriptions,           (MERGE, idempotent)  
     (schema_embeddings,  
      table descriptions,  
      column_embeddings,  
      example queries)  
     query_memory)  
2. BigQuery tables ....> ML.GENERATE_EMBEDDING ....> embedding column  
filled  
    (rows with empty             (text-embedding-005,  
     (768-dim FLOAT64 array)       RETRIEVAL_DOCUMENT)  
     embedding arrays)
```

Vector Search at Query Time

```
Trader question: "what was the average edge today?"  
|  
v  
ML.GENERATE_EMBEDDING(question, task_type='RETRIEVAL_QUERY')  
|  
v  
VECTOR_SEARCH(schema_embeddings, COSINE, top_k=5)  
|  
v  
Results: [{table: markettrade, dataset: kpi, distance: 0.12},  
...]
```

Key details:

- **RETRIEVAL_DOCUMENT** task type for stored content (when embedding table/column descriptions)
- **RETRIEVAL_QUERY** task type for search queries (when embedding the trader's question)
- **COSINE distance** (lower = more similar)
- **No vector indexes** needed (< 200 rows, brute-force search is fast)
- All operations are **idempotent** (CREATE OR REPLACE, MERGE)

Continuous Learning Loop

When a trader confirms a query result is correct, `save_validated_query` inserts the Q->SQL pair into `query_memory` with an embedding. Next time a similar question comes in, `fetch_few_shot_examples` retrieves it as a reference for the LLM.

```
Trader: "Is this what you were looking for?"
```

```
Trader: "Yes"
```

```
|  
v
```

```
save_validated_query(question, sql, tables_used, ...)
|
v
INSERT into query_memory + ML.GENERATE_EMBEDDING
|
v
Future similar questions find this pair via VECTOR_SEARCH
```

The 6-Tool Chain in Detail

1. vector_search_tables(question)

Finds which BigQuery tables are relevant to a question using semantic similarity against schema_embeddings.

Input: Natural language question **Output:** Top-5 tables ranked by relevance with descriptions

2. load_yaml_metadata(table_name, dataset_name)

Loads the full YAML catalog for a specific table -- column names, types, descriptions, synonyms, business rules. Also loads dataset context for KPI/data routing.

Input: Table name + dataset name (for disambiguation) **Output:** Full YAML metadata as string

3. fetch_few_shot_examples(question)

Retrieves similar past validated queries from query_memory via vector search. Provides proven SQL patterns for the LLM to follow.

Input: Natural language question **Output:** Up to 5 similar Q->SQL pairs with routing signals

4. dry_run_sql(sql_query)

Validates SQL syntax, column references, table permissions, and estimates cost -- all without executing.

Input: SQL query string **Output:** Valid/invalid status + estimated bytes/MB

5. execute_sql(sql_query)

Executes the validated query. Enforces read-only (SELECT/WITH only), auto-adds LIMIT 1000.

Input: SQL query string **Output:** Row count + result rows as dicts

6. save_validated_query(question, sql_query, ...)

Saves a confirmed-correct Q->SQL pair for future retrieval. Immediately embeds it.

Input: Question, SQL, tables used, complexity, routing signal **Output:** Success/error status

Safety Mechanisms

Before-Tool Guard

Every call to dry_run_sql or execute_sql passes through before_tool_guard:

- Extracts the first SQL keyword
- Allows only **SELECT** and **WITH** (CTE)
- Blocks **INSERT, UPDATE, DELETE, DROP, ALTER, CREATE**
- Returns an error dict that the LLM sees as a tool failure

System Prompt Enforcement

The system prompt explicitly instructs:

- NEVER generate INSERT, UPDATE, DELETE, DROP, CREATE, ALTER, or any DDL/DML
- NEVER query tables not listed in the catalog
- NEVER use SELECT * -- always specify columns
- ALWAYS filter on `trade_date` partition column
- ALWAYS use fully-qualified table names
- ALWAYS add LIMIT unless user asks for all rows

Read-Only Architecture

- BigQuery client uses ADC with read-only scopes
 - `execute_sql` double-checks first keyword before execution
 - No write credentials are configured
-

Conductor Tracks (Implementation Phases)

Track 01: Foundation -- COMPLETE

Established the project skeleton and infrastructure.

Delivered:

- Repository structure, Dockerfile, docker-compose.yml
- ADK agent skeleton (`root_agent + n12sql_agent`)
- Configuration system (pydantic-settings with .env)
- BigQuery client with Protocol-based dependency injection
- Structured JSON logging via structlog
- Dev GCP project setup with sample data

Track 02: Context Layer -- COMPLETE

Built the two-layer metadata system that gives the agent deep knowledge of the data.

Delivered:

- 15 YAML catalog files covering all 13 tables
- Dataset-level routing rules and disambiguation logic
- 30+ validated Q->SQL example pairs across 3 example files
- BigQuery embedding infrastructure (3 tables)
- Embedding pipeline scripts (idempotent, CREATE OR REPLACE / MERGE)
- Vector search validation (5/5 test queries returning correct tables)
- `_routing.yaml` with critical cross-dataset routing patterns

Track 03: Agent Tools -- COMPLETE

Implemented the 6-tool chain that the LLM uses to answer questions.

Delivered:

- `vector_search_tables` -- semantic table routing via VECTOR_SEARCH
- `fetch_few_shot_examples` -- retrieve similar validated queries
- `load_yaml_metadata` -- load YAML catalog for specific tables
- `dry_run_sql` -- validate SQL syntax and permissions
- `execute_sql` -- run query with read-only enforcement
- `save_validated_query` -- continuous learning loop
- Shared dependency injection (`_deps.py`)
- 115 unit tests

Track 04: Agent Logic -- COMPLETE

Wired the tools into the agent with a comprehensive system prompt and safety callbacks.

Delivered:

- Dynamic system prompt (`build_nl2sql_instruction`) with:
 - 7 routing rules, tool usage order, SQL generation rules
 - Date injection, clarification rules, response format
 - Full table catalog embedded in prompt
- `before_tool_guard` callback (blocks DML/DDL)
- `after_tool_log` callback (observability)
- `GenerateContentConfig(temperature=0.1)` for deterministic SQL
- 155 unit tests + 20 integration tests

Track 05: Eval & Hardening -- PLANNED

Will establish accuracy metrics and harden the agent.

Goals:

- Gold-standard evaluation set (50+ questions with expected SQL)
- Accuracy metrics per question type and table
- Retry logic for failed queries (up to 3 attempts)
- Edge case handling (empty results, ambiguous questions)

Track 06: Metadata Enrichment -- PLANNED

Will expand metadata coverage and improve routing accuracy.

Goals:

- Additional example queries per table
- Enriched column descriptions from actual data profiling
- Portfolio-specific routing rules
- Column-level vector search (using `column_embeddings`)

Environment Configuration

| Setting | Dev | Production ||-----|----|-----|| **GCP Project** | melodic-stone-437916-t3 |
cloud-data-n-base-d4b3 || **LiteLLM Proxy** | http://localhost:4000 |
<https://litellm.production.mako-cloud.com/> || **LLM Models** | openai/claude-haiku /
openai/claude-sonnet | openai/gemini-3-flash-preview /
openai/gemini-3-pro-preview || **Embedding Model** | Same project connection | Cross-project:
cloud-ai-d-base-a2df || **Data** | Sample/thin slices | Full production data |

Switching environments requires only changing the .env file -- no code changes.

Running the Agent

```
# Unit tests (default, no live services needed)
pytest tests/
# Integration tests (requires ADC + optionally LiteLLM proxy)
pytest -m integration tests/
# Docker web UI
docker compose up
# Open http://localhost:8000 -> select nl2sql_agent
# Docker terminal mode
docker compose run --rm agent adk run nl2sql_agent
```